

Bounds and Algorithms for Fast Exact Searches of Chemical Fingerprints in Linear and Sublinear Time

S. Joshua Swamidass and Pierre Baldi*

Institute for Genomics and Bioinformatics, School of Information and Computer Sciences,
University of California, Irvine, Irvine, California 92697-3435

Received August 16, 2006

Chemical fingerprints are used to represent chemical molecules by recording the presence or absence, or by counting the number of occurrences, of particular features or substructures, such as labeled paths in the 2D graph of bonds, of the corresponding molecule. These fingerprint vectors are used to search large databases of small molecules, currently containing millions of entries, using various similarity measures, such as the Tanimoto or Tversky's measures and their variants. Here, we derive simple bounds on these similarity measures and show how these bounds can be used to considerably reduce the subset of molecules that need to be searched. We consider both the case of single-molecule and multiple-molecule queries, as well as queries based on fixed similarity thresholds or aimed at retrieving the top K hits. We study the speedup as a function of query size and distribution, fingerprint length, similarity threshold, and database size $|D|$ and derive analytical formulas that are in excellent agreement with empirical values. The theoretical considerations and experiments show that this approach can provide linear speedups of one or more orders of magnitude in the case of searches with a fixed threshold, and achieve sublinear speedups in the range of $O(|D|^{0.6})$ for the top K hits in current large databases. This pruning approach yields subsecond search times across the 5 million compounds in the ChemDB database, without any loss of accuracy.

1. INTRODUCTION

One of the most fundamental tasks of chemoinformatics is the rapid search of large repositories of molecules. In a typical search, given a query consisting of a molecule or a family of molecules, one is interested in retrieving all the molecules contained in a large repository, such as PubChem, ZINC,¹ or ChemDB,² that are similar to the query and satisfy the given constraints. To facilitate this process, in many chemoinformatics systems, molecules are represented by binary fingerprint vectors (refs 3–7 and references therein). It is these fingerprints and their similarity measures that are used to search these large repositories. While fingerprint representations yield efficient search algorithms, it is still important to keep search times to a minimum to allow interactive searches and to allow the search space to grow considerably beyond its current typical value of a few million molecules toward the estimated 10^{60} size of the virtual space of small organic molecules.⁸ To further reduce search times, here we first derive bounds on all the standard similarity measures and then show how these bounds can be used to prune the search space and considerably speed up current searches, without any loss of accuracy.

2. MOLECULAR FINGERPRINT NOTATION AND SIMILARITY MEASURES

2.1. Molecular Fingerprint Notation. Let \mathcal{A} denote a molecule and $\vec{A} = (A_i)$ denote the corresponding fingerprint vector, with $1 \leq i \leq N$. The precise interpretation of the fingerprint components is irrelevant for our purpose. As in

most chemoinformatics systems, one may consider that each component A_i is a bit associated with the presence or absence of a particular substructure (e.g., functional group, labeled path, labeled cycle, or labeled tree) of atoms and bonds in the molecule. While many systems use binary fingerprints, it is also possible to use richer, integer-valued fingerprints where the components A_i count the number of occurrences of the corresponding substructure. Here, the length N of the fingerprints is not important—our considerations apply to full-length fingerprints (large values of N), compressed fingerprints (typically $N = 512$ or $N = 1024$), or variable-length fingerprints. For each fingerprint \vec{A} , we let $A = \sum_i A_i$. In the binary case, A is the total number of bits set to 1. In the general case, A is the sum of all substructure counts. In the binary case, we also use $A \cup B$ (respectively $A \cap B$) to denote the total number of bits set to 1 in \vec{A} OR \vec{B} (respectively \vec{A} AND \vec{B}).

While in current systems queries are often based on a single molecule, sometimes it is useful to use several related molecules in a query of the form $\vec{\mathcal{A}} = \{\mathcal{A}_1, \dots, \mathcal{A}_M\}$, as in the example of Figure 1. In this notation, \mathcal{A}_i represents the i th molecule in the family, and $M = |\mathcal{A}|$ represents the number of molecules in the query. In this case, $\vec{A}_i = (A_{ij})$ denotes the fingerprint vector associated with molecule \mathcal{A}_i , and $A_i = \sum_j A_{ij}$.

2.2. Molecular Similarity Measures. The measure used to assess similarity between molecules, hence, fingerprints, plays a fundamental role in chemical searches. Several similarity measures have been introduced for molecular fingerprints, two of the most common ones being the Tanimoto measure and its generalization, the Tversky measure, for binary fingerprints.^{9,10} Both measures can be

* Corresponding author e-mail: pfbaldi@ics.uci.edu.




Figure 1. Examples of five estrogen receptor binding compounds that could be entered in a multiple-molecule query aimed at retrieving additional compounds in the same family.

generalized to count-based fingerprints using a MinMax operator,¹¹ as described below. Our derivations here are illustrated primarily with the Tanimoto and Tversky measures in the binary case, since these are the most widely used, but we treat also their extensions to the nonbinary case. Moreover, the same ideas can be applied immediately to many other fingerprint similarity measures, as described in the appendices.

The Tanimoto similarity measure between two binary fingerprints is defined by the ratio of the number of common bits set to 1 to the total number of bits set to 1 in the two fingerprints

$$S(\mathcal{A}, \mathcal{B}) = S(\vec{A}, \vec{B}) = (A \cap B) / (A \cup B) \quad (1)$$

The Tanimoto coefficient is essentially a variation on the F measure of information retrieval.¹²

The Tversky measure with parameters α and β between two binary fingerprints is defined by

$$S_{\alpha\beta}(\vec{A}, \vec{B}) = \frac{A \cap B}{\alpha A + \beta B + (1 - \alpha - \beta)(A \cap B)} \quad (2)$$

with $0 \leq \alpha$ and $0 \leq \beta$. When $\alpha = \beta = 1$, it reduces to the Tanimoto measure. In the nonsymmetric case ($\alpha \neq \beta$), α and β are used for biasing the search toward superstructures or substructure of the query molecule \mathcal{A} . A large relative value of α will bias the search toward superstructures of \mathcal{A} , whereas a large relative value of β will bias the search toward substructures of \mathcal{A} .

For nonbinary fingerprints associated with actual counts, the MinMax measure^{11,13} is given by

$$S^C(\vec{A}, \vec{B}) = \frac{\sum_i \min(A_i, B_i)}{\sum_i \max(A_i, B_i)} \quad (3)$$

The MinMax measure reduces to the Tanimoto measure in the case of binary fingerprints.

Finally, for nonbinary fingerprints, we can also generalize the Tversky measure to get the MinMax Tversky measure

$$S_{\alpha\beta}^C(\vec{A}, \vec{B}) = \frac{\sum_i \min(A_i, B_i)}{\alpha \sum_i A_i + \beta \sum_i B_i + (1 - \alpha - \beta) \sum_i \min(A_i, B_i)} \quad (4)$$

which reduces to the Tversky measure in the case of binary fingerprints and to the MinMax measure in the symmetric case with $\alpha = \beta = 1$.

3. BOUNDS ON SIMILARITY FOR SINGLE-MOLECULE QUERY

Let \vec{A} be the query fingerprint. We can calculate efficient upper bounds for all the similarity measures described in the literature (see also the appendices). The upper bound for the Tanimoto similarity can be computed by writing the similarity as $S(\vec{A}, \vec{B}) = (A \cap B) / (A + B - A \cap B)$ and noticing that the derivative of $f(x) = x / (A + B - x)$ with respect to x is positive. Thus, the derivative of the Tanimoto similarity is positive with respect to $(A \cap B)$, and an upper bound can be derived by setting the number of bits in common $(A \cap B)$ to its maximum possible value: $\min(A, B)$. In other words, for fixed \vec{A} and \vec{B}

$$S(\vec{A}, \vec{B}) \leq T(A, B) = \frac{\min(A, B)}{A + B - \min(A, B)} = \frac{\min(A, B)}{\max(A, B)} \quad (5)$$

Here and everywhere else, S is used to denote the similarity measure and T the bound.

Using similar reasoning, we can derive a more general bound on the Tversky similarity in the form

$$S_{\alpha\beta}(\vec{A}, \vec{B}) \leq T_{\alpha\beta}(A, B) = \frac{\min(A, B)}{\alpha A + \beta B + (1 - \alpha - \beta) \min(A, B)} \quad (6)$$

In the case of the MinMax similarity for nonbinary vectors, we have

$$S^C(\vec{A}, \vec{B}) \leq T^C(A, B) = \frac{\min(A, B)}{A + B - \min(A, B)} = \frac{\min(A, B)}{\max(A, B)} \quad (7)$$

This bound uses the fact that $\sum_i \min(A_i, B_i) \leq \min(A, B)$ and $\sum_i \max(A_i, B_i) = A + B - \sum_i \min(A_i, B_i)$. It is also a special case of the bound on the MinMax Tversky similarity with $\alpha = \beta = 1$, as given below.

In the case of the MinMax Tversky similarity for nonbinary vectors, by combining the arguments above, we get the bound

$$S_{\alpha\beta}^C(\vec{A}, \vec{B}) \leq T_{\alpha\beta}^C(A, B) = \frac{\min(A, B)}{\alpha A + \beta B + (1 - \alpha - \beta) \min(A, B)} \quad (8)$$

It is worth noting that all the general bounds we have derived are optimal in the sense that one can construct examples where these bounds are actually achieved.

4. BOUNDS ON SIMILARITY FOR MULTIPLE-MOLECULE QUERY

When the query is a family of molecules $\vec{\mathcal{A}} = \{\mathcal{A}_1, \dots, \mathcal{A}_M\}$, there are two basic ways of defining a

similarity measure $S(\overline{\mathcal{A}}, \overline{\mathcal{B}}) = S(\mathcal{A}_1, \dots, \mathcal{A}_M, \mathcal{B}) = S(\overline{A}_1, \dots, \overline{A}_M, \overline{B})$ between the query family and each molecule \mathcal{B} in the database: (1) by aggregating the individual pairwise similarity measures $S(\mathcal{A}_i, \mathcal{B}) = S(A_i, B)$ and (2) by aggregating the fingerprints \overline{A}_i into a profile fingerprint \overline{P} and then using a similarity measure of the form $S(\overline{P}, \overline{B})$. With obvious adjustments, the ideas presented here can easily be extended to “family–family” (or “profile–profile”) comparisons of the form $S(\overline{\mathcal{A}}, \overline{\mathcal{B}}) = S(\mathcal{A}_1, \dots, \mathcal{A}_M, \mathcal{B}_1, \dots, \mathcal{B}_Q)$ between two families $\overline{\mathcal{A}}$ and $\overline{\mathcal{B}}$ of molecules, where $\overline{\mathcal{B}} = \{\mathcal{B}_1, \dots, \mathcal{B}_Q\}$.

4.1. Bounds on Aggregated Similarity for Multiple-Molecule Query. In the first class of approaches, the similarity $S(\overline{\mathcal{A}}, \overline{\mathcal{B}})$ is defined in terms of the individual similarities $S(\mathcal{A}_i, \mathcal{B})$, for instance, by taking their weighted average, maximum, or minimum. In each one of these cases, bounds can easily be derived using the bounds obtained in section 3. In the case of a weighted average with non-negative weights $w_1 \dots w_M$, we have

$$S(\mathcal{A}_1, \dots, \mathcal{A}_M, \mathcal{B}) = \sum_i w_i S(\overline{A}_i, \overline{B}) \leq \sum_i w_i T(A_i, B) \quad (9)$$

where S can be any of the similarities defined above and T the corresponding bound. Likewise, in the case where the family similarity is defined by the minimum pairwise similarity, we have

$$S(\mathcal{A}_1, \dots, \mathcal{A}_M, \mathcal{B}) = \min_i S(\overline{A}_i, \overline{B}) \leq \min_i T(A_i, B) \quad (10)$$

and for the maximum

$$S(\mathcal{A}_1, \dots, \mathcal{A}_M, \mathcal{B}) = \max_i S(\overline{A}_i, \overline{B}) \leq \max_i T(A_i, B) \quad (11)$$

Within the first class of approaches, it is also possible to consider measures that are obtained by combining the results of elementary pairwise comparisons between the molecules (e.g., intersections and unions) rather than the similarity measures themselves. In particular, one can derive a series of measures by simply aggregating the numerators and the denominators of the previous four similarity measures of section 2. For instance, in the case of the binary Tversky measure, we can define an aggregate similarity measure:

$$S_{\alpha\beta\overline{w}}(\mathcal{A}_1, \dots, \mathcal{A}_M, \mathcal{B}) = \frac{\sum_i w_i (A_i \cap B)}{\sum_i w_i [\alpha_i A_i + \beta_i B + (1 - \alpha_i - \beta_i)(A_i \cap B)]} \quad (12)$$

Note that each molecule \mathcal{A}_i comes with two sets of parameters: a weight w_i , which measures its importance among the family, and Tversky's parameters α_i and β_i , which can be used to bias the search toward substructures or superstructures of \mathcal{A}_i . A special case of this particular measure is used in Xue et al.⁷

We can derive bounds on this measure by noticing again that the similarity has a positive derivative with respect to each term $A_i \cap B$ and using the inequality $A_i \cap B \leq \min_i(A_i, B)$. This yields

$$S_{\alpha\beta\overline{w}}(\mathcal{A}_1, \dots, \mathcal{A}_M, \mathcal{B}) \leq \frac{T_{\alpha\beta\overline{w}}(A_1, \dots, A_M, B)}{\sum_i w_i [\alpha_i A_i + \beta_i B + (1 - \alpha_i - \beta_i) \min(A_i, B)]} \quad (13)$$

Likewise, a weighted MinMax measure can be written as

$$S_w^C(\mathcal{A}_1, \dots, \mathcal{A}_M, \mathcal{B}) = \frac{\sum_i w_i \sum_j \min(A_{ij}, B_j)}{\sum_i w_i \sum_j \max(A_{ij}, B_j)} \quad (14)$$

Here, A_{ij} denotes the j th component of the fingerprint \overline{A}_i .

With even greater generality, we can consider a weighted MinMax Tversky similarity of the form

$$S_{\alpha\beta\overline{w}}^C(\mathcal{A}_1, \dots, \mathcal{A}_M, \mathcal{B}) = \frac{\sum_i w_i \sum_j \min(A_{ij}, B_j)}{\sum_i w_i [\alpha_i A_i + \beta_i B + \sum_j (1 - \alpha_i - \beta_i) \min(A_{ij}, B_j)]} \quad (15)$$

with the generalized bounds

$$S_{\alpha\beta\overline{w}}^C(\mathcal{A}_1, \dots, \mathcal{A}_M, \mathcal{B}) \leq \frac{T_{\alpha\beta\overline{w}}^C(A_1, \dots, A_M, B)}{\sum_i w_i [\alpha_i A_i + \beta_i B + (1 - \alpha_i - \beta_i) \min(A_i, B)]} \quad (16)$$

which can again be shown using some algebra and noticing that the derivative of eq 15 is positive with respect to $\sum_j \min(A_{ij}, B_j)$ and that $\sum_j \min(A_{ij}, B_j) \leq \min(A_i, B)$.

4.2. Bounds on Profile Similarity for Multiple-Molecule Query. A second possible class of approaches to multiple-molecule queries is to build a fingerprint profile $\overline{P}(\overline{\mathcal{A}}) = \overline{P} = (P_j)$ to represent the family $\overline{\mathcal{A}}$ and then measure similarity between the profile vector and the fingerprint vectors of the molecules to be searched. A fingerprint profile summarizes the information in a set of fingerprints, very much like a sequence profile or a position-specific scoring matrix summarizes the information in a set of aligned sequences in bioinformatics. Fingerprint profiles can in turn be subdivided into linear and nonlinear profiles.

A linear fingerprint profile stores for each component the frequency that the corresponding bit is set to 1 in the family of fingerprints. If a given bit position is set to 1 in half of the fingerprints in the family, then the corresponding component in the profile is set to 0.5. If the fingerprint consists of integer counts rather than bits, then the profile consists of average counts. In addition, different weights can also be assigned to each molecule or fingerprint in the family. In this more general case of linear profiles, $P_j^w = \sum_i w_i A_{ij}$. For proper scaling, it is desirable to use a convex combination with $w_i \geq 0$ and $\sum_i w_i = 1$. Throughout this section, $i = 1, \dots, M$ runs over the molecules in the family, and $j = 1, \dots, N$ runs over the fingerprint components.

The similarity between the profile $\overline{P}(\overline{\mathcal{A}})$ and a fingerprint \overline{B} can be measured using the MinMax measure

$$S^C(\overline{P}^w, \overline{B}) = \frac{\sum_j \min(P_j^w, B_j)}{\sum_j \max(P_j^w, B_j)} \quad (17)$$

or the generalized MinMax Tversky measure

$$S_{\alpha\beta}^C(\overline{P}^w, \overline{B}) = \frac{\sum_j \min(P_j^w, B_j)}{\alpha \sum_j P_j^w + \beta \sum_j B_j + (1 - \alpha - \beta) \sum_j \min(P_j^w, B_j)} \quad (18)$$

Letting $A^+ = P^w = \sum_j P_j^w = \sum_j \sum_i w_i A_{ij}$, we can use the bounds on these measures to derive the bound

$$S^C(\vec{P}^w, \vec{B}) \leq T^C(P^w, B) = \frac{\min(A^+, B)}{A^+ + B - \min(A^+, B)} \quad (19)$$

and its more general version

$$S_{\alpha\beta}^C(\vec{P}^w, \vec{B}) \leq T_{\alpha\beta}^C(P^w, B) = \frac{\min(A^+, B)}{\alpha A^+ + \beta B + (1 - \alpha - \beta) \min(A^+, B)} \quad (20)$$

In the case of a convex combination ($\sum_i w_i = 1$), we can derive an even better bound in the form

$$S_{\alpha\beta}^C(\vec{P}^w, \vec{B}) \leq T_{\alpha\beta}^C(A_1, \dots, A_M, B) = \frac{\sum_i w_i \min(A_i, B)}{\sum_i w_i [\alpha A_i + \beta B + (1 - \alpha - \beta) \min(A_i, B)]} \quad (21)$$

In other words, $S_{\alpha\beta}^C(\vec{P}^w, \vec{B}) \leq T_{\alpha\beta}^C(A_1, \dots, A_M, B) \leq T_{\alpha\beta}^C(P^w, B)$.

If the A_i 's are binary and the combination is convex, then we have the identity $\sum_j \min(P_j^w, B_j) = \sum_i w_i (A_i \cap B)$. Therefore, in the binary case with convex linear combination, we have the identities

$$S^C(\vec{P}^w, \vec{B}) = \frac{\sum_j \min(P_j^w, B_j)}{\sum_j \max(P_j^w, B_j)} = \frac{\sum_i w_i (A_i \cap B)}{\sum_i w_i [A_i + B - (A_i \cap B)]} \quad (22)$$

and

$$\begin{aligned} S_{\alpha\beta}^C(\vec{P}^w, \vec{B}) &= \frac{\sum_j \min(P_j^w, B_j)}{\sum_j \alpha P_j^w + \beta B_j + (1 - \alpha - \beta) \sum_j \min(P_j^w, B_j)} \\ &= \frac{\sum_i w_i (A_i \cap B)}{\sum_i w_i [\alpha A_i + \beta B + (1 - \alpha - \beta) (A_i \cap B)]} \end{aligned} \quad (23)$$

Thus, in the binary case, the MinMax and Tversky measures applied to the profile vectors \vec{P}^w and \vec{B} are equivalent to summing the numerators and denominators of the individual Tanimoto and Tversky measures between the A_i 's and B (eq 12), and thus the corresponding bound (eq 13) can be applied. Note that, in the binary case, since the second family of approaches (profile) is a special case of the first family of approaches (aggregation), for bound purposes it is *not* necessary to compute the actual profile vector \vec{P}^w . This is not true in the nonbinary case, or in the case of nonlinear profiles described below.

To build nonlinear profiles, a nonlinear transformation is applied to the frequency counts in each column in order to derive a profile vector. In the case of consensus fingerprints,^{6,15–17} for instance, the profile vector contains a 1 at a given position if and only if all the fingerprint vectors of the molecules in the query family also contain a 1 in that position. This can be generalized to modal fingerprints with threshold t ,^{16,17} by setting a bit to 1 in the profile vector if and only if the

proportion of molecules in the query family that have a 1 at the corresponding position is greater than or equal to t . Consensus profiles correspond to modal profiles with $t = 1$. Even more generally, a nonlinear function such as a logistic function can be used to map frequencies to profile components. In the case of nonlinear profiles, all the similarity bounds can still be used with the profile vector. The only difference with the linear case is that the bound in the nonlinear case may not be easily expressed in terms of the bounds derived from the individual molecules in the query family.

5. DATA

Before we show how the bounds can be used to speed up searches in large databases of compounds, we describe the data used in the experiments. All the data are extracted from the ChemDB database,² which currently contains on the order of 5 million (5M) unique compounds. In the experiments, we use fingerprints associated with labeled paths of length up to 8 (i.e., nine atoms and eight bonds). In this case, the total number of observed labeled paths is $N_* = 152\,087$. Compression is done using a simple modulo operator. Most results are reported for fingerprints of length $N = 512$. However, we have tested all values $N = 2^n$, with $5 \leq n \leq 10$, and report the corresponding results when the dependence on fingerprint length is relevant. Robust results are obtained by increasing the path length or varying N_* or N . All fingerprints are computed using an in-house program written in Python.

For experiments that require computing pairwise similarities between all molecules, we use a random data set of 50 000 molecules extracted from ChemDB, corresponding approximately to 1.25×10^9 pairwise similarity measurements. Varying the random sample does not affect these results in any significant way.

Figure 2 reports the distribution of A across the entire ChemDB database, together with the distribution of A over the queries received by ChemDB over the Web during a 4-month period (February 2006 to June 2006). The distribution of A across the entire database is well approximated by a Gaussian distribution with a mean of 119.53 and standard deviation of 40.07. In contrast, the mean and standard deviation of the actual queries are 64.09 and 45.88, respectively.

For some experiments, we use the Stahl and Rarey¹⁴ data sets which consist of six groups of diverse molecules with similar activity. The molecules of each group are known to interact with the same protein. These data sets consist of 128 chemicals which interact with Cox-2, 55 which interact with estrogen receptor, 43 which interact with gelatinase-A, 17 which interact with neuraminidase, 25 which interact with p38-MAP kinase, and 67 which interact with thrombin. All data sets are available upon request.

6. RESULTS: FAST SEARCH ALGORITHMS

6.1. Fast Search for Single-Molecule Query. We can now show how the bounds derived in the previous sections can be used to accelerate searches. For a given query and similarity threshold, we only need to sift through a small subset of molecules: those that satisfy the corresponding bound. More precisely, for any similarity measure, if we are




Figure 2. Empirical distributions of $A = \sum_i A_i$ over the entire ChemDB, and over the set of ChemDB queries received over the Web during a period of four consecutive months, using fingerprints of length $N = 512$. The distribution of A across the entire database is well-modeled using a Gaussian distribution with a mean of 119.53 and standard deviation of 40.07. In contrast, the mean and standard deviation of the received queries are 64.09 and 45.88, respectively.




Figure 3. Pruning the search space. The binary fingerprint of the query molecule satisfies $A = 100$. As a function of B , the Tanimoto similarity measure $S(A,B)$ is upper-bounded by the curve $T(A,B)$. If the similarity threshold is set at 0.9, only molecules with B in a very small interval around 100 need to be searched. All other molecules have similarity scores that are below the threshold.

interested in retrieving only molecules that have similarity to the query \mathcal{A} above a given threshold t ($0 \leq t \leq 1$), then we can discard all the molecules \mathcal{B} which satisfy $T(A,B) \leq t$. This can drastically reduce the number of molecules which must be examined, thereby improving speed without affecting accuracy at all.

For illustration purposes, consider first the simple case of binary fingerprints with the Tanimoto similarity measure. For a query fingerprint of size A , eq 5 shows that, if $B \leq A$,

then $T(A,B) = B/A$ is linear in B . If $B \geq A$, then $T(A,B) = A/B$ and decays like $1/B$ (Figure 3). Thus, all molecules with $(B/A) < t$ can be discarded. Likewise, all molecules with $(A/B) < t$ can be discarded. Thus, the search can be restricted to molecules with B satisfying

$$At \leq B \leq \frac{A}{t} \quad (24)$$




Figure 4. Bounds T on the Tanimoto measure as a function of the size B of the binary fingerprint of a molecule in the database. The solid curve corresponds to a single-molecule query with $A = 100$. The dashed curve corresponds to a two-molecule query with the two molecules satisfying $A = 100$ and $A = 150$, using the similarity measure in eq 12, with $\alpha_i = \beta_i = 1$ and $w_i = 0.5$ (also equivalent in this binary case to eq 17).




Figure 5. Bounds T on similarity measures as a function of the size B of the binary fingerprint of a molecule in the database. Two-molecule query with $A_1 = 100$ and $A_2 = 150$. Each curve corresponds to a different similarity measure. Average (green) is the average Tanimoto similarity across the two molecules in the query. Min (red dashed) [respectively Max (red)] is the minimum (respectively maximum) of the Tanimoto similarities. Aggregate (blue) is the aggregation of the individual Tanimoto similarity measures (eq 12 with $\alpha_i = \beta_i = 1$ and $w_i = 0.5$ for all i).

Likewise, for the Tversky similarity measure applied to binary fingerprints, we can discard all molecules with fingerprint vector B satisfying $T_{\alpha\beta}(A, B) < t$. A rearrangement of eq 6 shows that only molecules with fingerprint

vector \vec{B} satisfying

$$\frac{At\alpha}{1-t+t\alpha} \leq B \leq \frac{A(1-t+t\beta)}{t\beta} \quad (25)$$




Figure 6. Bounds on similarity measure. Query consists of 55 estrogen receptor binding molecules against the same random sample of 50 000 molecules extracted from the ChemDB database. Molecules are represented by binary compressed fingerprints of length 512. Each curve corresponds to the bound for a different similarity measure. Average (green) is for the average Tanimoto similarity across the 55 molecules in the query. Min (red dashed) [respectively Max (red)] is for the minimum (respectively maximum) of the Tanimoto similarities. Profile (blue) is for the aggregation of the 55 individual Tanimoto similarity measures (eq 12 with $\alpha_i = \beta_i = 1$ and $w_i = 1/55$ for all i).

need to be included in the search. Needless to say, the B values need to be computed only once for each molecule and then stored in the database. Therefore, these bounds and the corresponding search restrictions can be implemented very efficiently. Similar considerations hold for all the other similarity measures (see also the appendices).

6.2. Fast Search for Multiple-Molecule Query. The same approach can be used for multiple-molecule queries. Figure 4, for instance, shows two curves corresponding to the bounds T on the Tanimoto measure, with a single-molecule query satisfying $A = 100$, as well as a two-molecule query with the two molecules satisfying $A_1 = 100$ and $A_2 = 150$. Figure 5 displays the bounds on different aggregate measures for the same two query molecules. By the same argument used in the case of single-molecule queries, at any given threshold, we can draw a corresponding horizontal line and restrict the search to only those values of B for which the upper bound exceeds the line. For instance, in Figure 4 with a query consisting of two molecules satisfying $A_1 = 100$ and $A_2 = 150$ and a similarity threshold $t = 0.85$, we see that no molecule in the database can satisfy such conditions, and therefore, the result ought to be set to “empty” immediately.

A more comprehensive and specific example is given in Figure 6 where the query molecules consist of a set of 55 estrogen receptor binding compounds,¹⁴ five of which are depicted in Figure 1.

6.3. Top K Hits. Most often, however, one does not prespecify a similarity threshold. Rather, as in most search engines and information retrieval systems, one is interested in finding the list of the most relevant hits. The methods described above can easily be adapted to the more realistic

case where one is interested in retrieving the top K hits, without setting any arbitrary threshold. For this, we note that for single molecule queries the upper bound is unimodal. Thus, we can search for the top K hits by starting from the maximum (where $A = B$) and exploring discrete possible values of B right and left of the maximum. More precisely, for binary fingerprints, we first index the molecules in the database by their fingerprint bit count to enable efficient referencing of a particular bit count bin. Next, with respect to a particular query, we calculate the bound on the similarity for every bit count in the database. Then, we sort these bit counts by their associated bound and iterate over the molecules in the database, in order of decreasing bound. As we iterate, we calculate the similarity between the query and the database molecule and use a heap to efficiently track the top hits. The algorithm terminates when the lowest similarity value in the heap is greater than the bound associated with the current database bin. The algorithm is given in Chart 1 in simple pseudo-code form (Algorithm 1).

This algorithm is equally valid in the case of multiple-molecule queries. Even though the bounds for multiple-molecule queries can have multiple maxima (e.g., Figures 5 and 6), the bins are sorted in decreasing order of their bounds. Thus, bounds with either single or multiple maxima are handled uniformly.

7. RESULTS: SPEEDUP

If f is the fraction of the database discarded by using the bounds, $1 - f$ is the fraction of the database to be searched and $1/(1 - f)$ is the speedup factor. The exact fraction of

Chart 1

```

Pseudocode for Top K Search
Require: database of fingerprints  $\vec{B}$  binned by bit count  $B_s$ 
Ensure: hits contains top  $K$  hits which satisfy  $\text{SIMILARITY}(\vec{A}, \vec{B}) > T$ 
1:  $hits \leftarrow \text{MINHEAP}()$ 
2:  $bounds \leftarrow \text{LIST}()$ 
3: for all  $B$  in database do //iterate over bins
4:    $tuple \leftarrow \text{TUPLE}(\text{BOUND}(A,B), B)$ 
5:    $\text{LISTAPPEND}(bounds, tuple)$ 
6: end for
7:  $\text{QUICKSORT}(bounds)$  //NOTE: the length of  $bounds$  is bounded by  $N$ 
8: for all  $bound, B$  in  $bounds$  do //iterate in order of decreasing  $bound$ 
9:   if  $bound < T$  then
10:    break //threshold stopping condition
11:   end if
12:   if  $K \leq \text{HEAPSIZE}(hits)$  and  $bound < \text{MINSIMILARITY}(hits)$  then
13:     break //top-K stopping condition
14:   end if
15:   for all  $\vec{B}$  in  $\text{database}[B]$  do
16:      $S = \text{SIMILARITY}(\vec{A}, \vec{B})$ 
17:      $tuple \leftarrow \text{TUPLE}(S, \vec{B})$ 
18:     if  $S \leq T$  then
19:       continue //ignore this  $\vec{B}$  and continue to next
20:     else if  $\text{LENGTH}(hits) < K$  then
21:        $\text{HEAPPUSH}(hits, tuple)$ 
22:     else if  $S > \text{MINSIMILARITY}(hits)$  then
23:        $\text{HEAPPUSH}(hits, tuple)$ 
24:        $\text{HEAPPUSH}(hits, tuple)$ 
25:     end if
26:   end for
27: end for
28: return  $hits$ 

```

the database to be discarded from a given search and, hence, the speedup factor depends on many variables, primarily the query molecule(s), the similarity measure, and the similarity threshold t . In this section, we study f as a function of these variables, both empirically and analytically. When necessary, examples and details are given using the Tanimoto similarity measure, but extensions to other similarity measures should be obvious.

Theoretically, and in the experiments to be described, f alone provides a good assessment of the speedup because the overhead associated with the bounds is negligible. Evaluating the bounds for all the bin sizes takes at most $O(N)$ steps, where N is the fingerprint length. Sorting the bins by their associated bounds takes at most $O(N \log N)$ steps. In practice, these overhead values are entirely negligible because N is quite small, in comparison with the size $|D|$ of the database. Similar considerations apply for the top K hits search algorithm.

Using continuous notation, we let $g_D(A)$ be the continuous density approximation to the values of A across the database D . When necessary, we use g_Q to denote the density over the queries, which can be different from g_D (Figure 2). The density g_D is well-approximated using a Gaussian distribution with mean μ and standard deviation σ . These parameters in turn can depend on fingerprint length.

7.1. Speedup for Single-Molecule Query as a Function of Threshold and A . For a single-molecule query \mathcal{A} and a threshold t , the bounds we have derived show that we need to search at most all the molecules satisfying $x_0 \leq B \leq x_1$, where the values of x_0 and x_1 depend on the value of A for the query molecule and on the similarity measure. Thus, the fraction $1 - f$ of the database to be searched satisfies

$$1 - f(A) = \int_{x_0}^{x_1} g_D(B) dB = G_D(x_1) - G_D(x_0) \quad (26)$$

where G_D is the corresponding distribution function. With Tanimoto similarity, for instance, combining eqs 24 and 26 with a Gaussian approximation yields

$$1 - f(A) = \int_{At}^{At/t} g_D(B) dB = F\left(\frac{\left(\frac{A}{t}\right) - \mu}{\sigma}\right) - F\left(\frac{At - \mu}{\sigma}\right) \quad (27)$$

where F is the distribution function of the normalized standard Gaussian distribution. This expression can of course be computed numerically. When t is very close to 1, we can use the approximation

$$1 - f(A) \approx \frac{A(1 - t^2)}{t\sigma} \frac{1}{\sqrt{2\pi}} e^{-(A-\mu)^2/2\sigma^2} \approx (1 - t) \frac{2A}{\sqrt{2\pi}\sigma} e^{-(A-\mu)^2/2\sigma^2} \quad (28)$$

so the speedup factor scales at least like $C/(1 - t)$ when t is close to 1.

Figure 7 provides contour plots showing the fraction of the database $f(A)$ to be excluded from a given search for the Tanimoto measure applied to compressed binary fingerprints of size 512, as a function of the query fingerprint size A , and for different threshold values: $t = 0.7, 0.8$, and 0.9 . The fundamental result is that this fraction is very significant and varies roughly from 30% to 100%, yielding searches that are faster by 1 or more orders of magnitude. In the worst case scenario, corresponding to $A \approx 120$ for $N = 512$ or $A^* \approx 139$ in the uncompressed fingerprint, the fraction varies from about 30% at $t = 0.7$ to about 75% at $t = 0.9$. In more favorable scenarios where A tends to be small or large, then the fraction is even larger. Note that for a fixed threshold the speedup factor is constant as a function of database size.

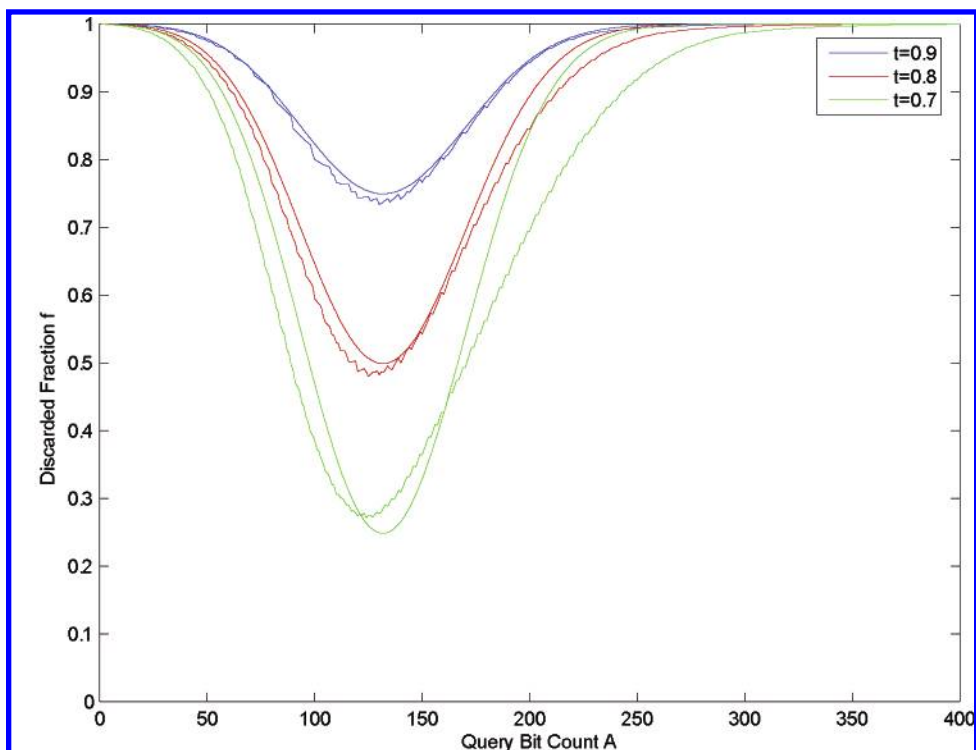


Figure 7. Curves representing the fraction f of molecules discarded from a given search as a function of the size A of the binary query fingerprint and the threshold t on the Tanimoto similarity. Blue corresponds to $t = 0.9$, red to $t = 0.8$, and green to $t = 0.7$. Results computed using a random sample of 50 000 molecules from the ChemDB database, using binary compressed fingerprints of length 512. The rough lines are the empirical curves. The smooth lines in this plot are the predictions provided by eq 28.

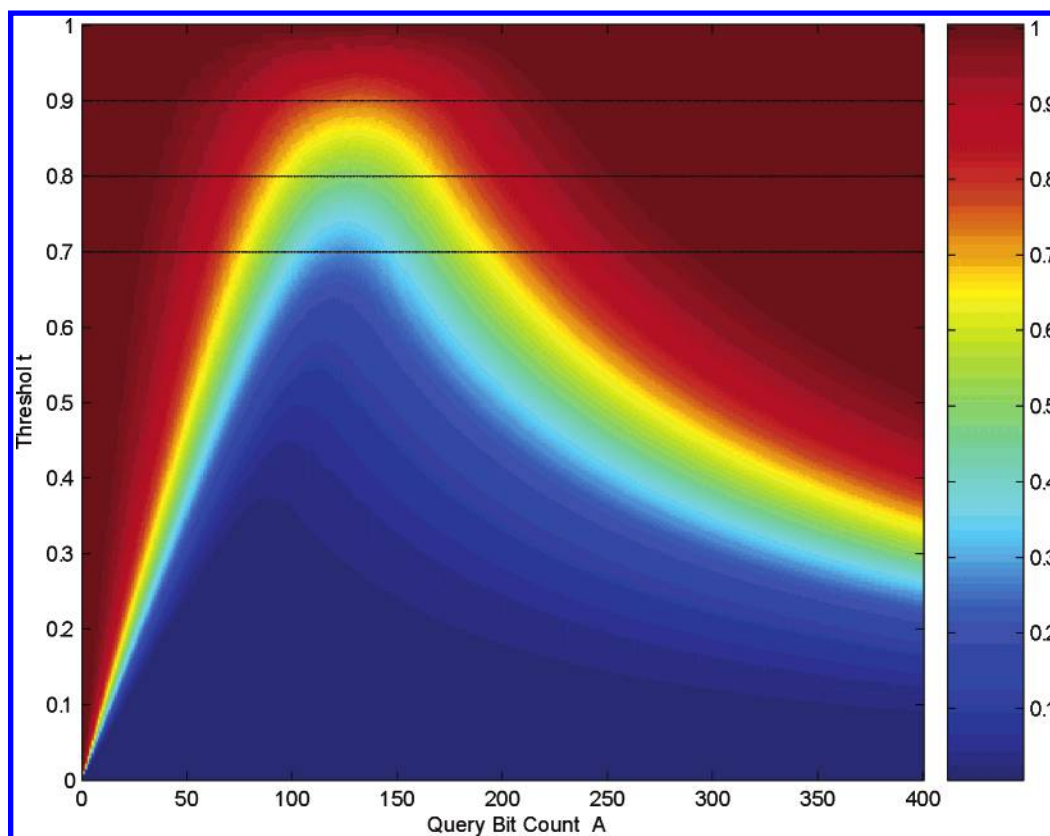


Figure 8. Discarded fraction f as a function of the query bit count A and the threshold t . The three lines associated with thresholds 0.7, 0.8, and 0.9 correspond to the curves in Figure 7.

As can be seen, the empirical curves agree very well with the theoretical approximation derived in eq 28. Figure 8 provides a more complete picture of f as a function of A and t .

7.2. Average Speedup for Single-Molecule Query as a Function of Threshold, A , and Fingerprint Length. Perhaps more important than the previous worst case analysis is the analysis of the average speedup. We can compute the




Figure 9. Average speedup as a function of the threshold for the single-molecule query using the Tanimoto similarity measure on binary compressed fingerprints of length 512. The average is computed over the ChemDB distribution using a random sample of 50 000 molecules (solid line), or over the distribution of Web queries received by ChemDB over a period of four consecutive months (dashed line). A similarity threshold of 0.9 yields approximately a 12-fold speedup for the average query.

average fraction of the database to be discarded by integrating over the query molecules

$$f = \int_D f(A) g_D(A) dA \quad (29)$$

This average can be computed with D equal to the entire database and the corresponding density $g_D(B)$, or using a density $g_Q(B)$ over queries. The latter in general gives an even more favorable speedup factor because queries tend to have a skewed distribution with respect to the database distribution, often with a smaller average value of A (Figure 9). For the average computed with g_D over the entire database, with the Tanimoto similarity measure, we can again use the Gaussian approximation to g_D to obtain

$$1 - f_D \approx \int_{A \in D} \left[G_D\left(\frac{A}{t}\right) - G_D(At) \right] g_D(A) dA \quad (30)$$

which can be estimated numerically. When t is close to 1, we can use the approximation in eq 28 and integrate to get

$$1 - f_D \approx \frac{\mu(1 - t^2)}{2t\sigma\sqrt{\pi}} \approx \frac{\mu(1 - t)}{\sigma\sqrt{\pi}} \quad (31)$$

thus, in this case, the speedup factor grows again like $C/(1 - t)$. Figure 9 shows how the average discarded fraction f is even larger when the average is computed over a distribution of typical queries (g_Q), rather than the database distribution (g_D).

The same techniques can be applied to study other effects, such as the impact of fingerprint length on the speedup factor. Figure 10 shows how the speedup increases with the length of the fingerprints. The fingerprint bit count distribution is more concentrated for shorter fingerprints, and therefore the

fraction of the database discarded at a given threshold is smaller. To obtain an analytical expression, we have derived in the appendices the following average relationship between A^* and A

$$A \approx N \left[1 - \left(1 - \frac{A^*}{N^*} \right)^{N^*/N} \right] \quad (32)$$

where A^* denotes the number of bits set to 1 in the unfolded (uncompressed) fingerprint, and N^* denotes the length of the uncompressed fingerprints. Combined with the Gaussian approximation to the distribution of A^* , this yields the distribution $g_D(A)$ as a function of the length N of the fingerprints

$$g_D(A) \approx \frac{N}{N - A} g_{N^*} \left[-N \log \left(1 - \frac{A}{N} \right) \right] \quad (33)$$

where g_{N^*} is the Gaussian approximation to the distribution over the uncompressed fingerprints of length N^* , with mean μ^* and standard deviation σ^* . Thus, in this case, we have

$$1 - f(A) = \int_{x_0}^{x_1} \frac{N}{N - A} g_{N^*} \left[-N \log \left(1 - \frac{A}{N} \right) \right] dA \quad (34)$$

With the Tanimoto measure, when t is close to 1, we get the approximation

$$1 - f(A) \approx \frac{A(1 - t^2)}{t} \frac{N}{N - A} g_{N^*} \left[-N \log \left(1 - \frac{A}{N} \right) \right] \approx \frac{2AN(1 - t)}{N - A} \frac{1}{\sqrt{2\pi}\sigma^*} e^{-\{\mu^* + N \log[1 - (A/N)]\}^2 / 2\sigma^{*2}} \quad (35)$$




Figure 10. Average discarded fraction as a function of the threshold for the single-molecule query using the Tanimoto similarity measure on binary compressed fingerprints of various lengths. The fraction increases monotonically with the length of the fingerprints.

By averaging over the database distribution (eq 30), the average fraction of the database to be searched with the Tanimoto similarity measure and similarity threshold t is given by

$$1 - f_D \approx \int_{A \in D} \left[G_D\left(\frac{A}{t}\right) - G_D(At) \right] \frac{N}{N-A} g_{N_s} \left[-N \log\left(1 - \frac{A}{N}\right) \right] dA \quad (36)$$

This can be computed by using the empirical or approximate distribution G_D for fingerprints of length N . It can also be approximated using again the approximate monotonic relation between A and A^* (eq 32) to yield

$$\int_{A \in D} \left\{ G_{N_s} \left[-N \log\left(1 - \frac{A/t}{N}\right) \right] - G_{N_s} \left[-N \log\left(1 - \frac{At}{N}\right) \right] \right\} \frac{N}{N-A} g_{N_s} \left[-N \log\left(1 - \frac{A}{N}\right) \right] dA \quad (37)$$

When t is close to 1, an analytic approximation can be derived as above. Table 1 shows that this formula is in very good agreement with empirical values, with less than 2% error at $t = 0.9$, for values of N ranging from 64 to 1024.

7.3. Speedup for the Top K Hits. For fixed threshold searches, the speedup does not depend on the size $|D|$ of the database. The situation is different, however, if we search for the top K hits. Using continuous notation, here, we let $f_A(x)$ and $F_A(x)$ respectively denote the density and distribution functions of the similarity scores (for instance Tanimoto scores) across the entire database averaged over all single-query molecules \mathcal{A} satisfying the constraint $\sum_i A_i = A$. Likewise, we can consider the average f_D of f_A , computed over the entire database, and the corresponding distribution F_D . The densities f_A and f_D are typically bell-shaped over

Table 1. Comparison between Predicted and Empirical Values of the Curves in Figure 10 as a Function of Fingerprint Length When the Threshold Approaches 1 (here $t = 0.9$)^a

N	empirical f	predicted value (eq 37)	percent error
1024	0.8345	0.835	0.06%
512	0.8226	0.8223	0.04%
256	0.7949	0.793	0.24%
128	0.7236	0.7148	1.22%
64	0.5057	0.4979	1.54%

^a Predicted values are derived from eq 37. The distribution g_{N_s} is approximated with a Gaussian with mean = 138.68 and standard deviation of 53.24, estimated from ChemDB.

the finite $[0,1]$ interval and therefore, to a first order of approximation, can be modeled using a β distribution of the form

$$\gamma x^{a-1} (1-x)^{b-1} = \frac{\Gamma(a+b)}{\Gamma(a)\Gamma(b)} x^{a-1} (1-x)^{b-1} \quad (38)$$

where a and $b \geq 0$ are parameters that depend on A in the case of f_A , and γ is the normalizing constant. With a database of size $|D|$ and a query of bit size A , the top K hits correspond to a threshold value u such that

$$\int_u^1 \gamma x^{a-1} (1-x)^{b-1} = \frac{K}{|D|} \quad (39)$$

By bounding and integrating, this gives immediately

$$\gamma u^{a-1} \frac{(1-u)^b}{b} \leq \frac{K}{|D|} \leq \gamma \frac{(1-u)^b}{b} \quad (40)$$

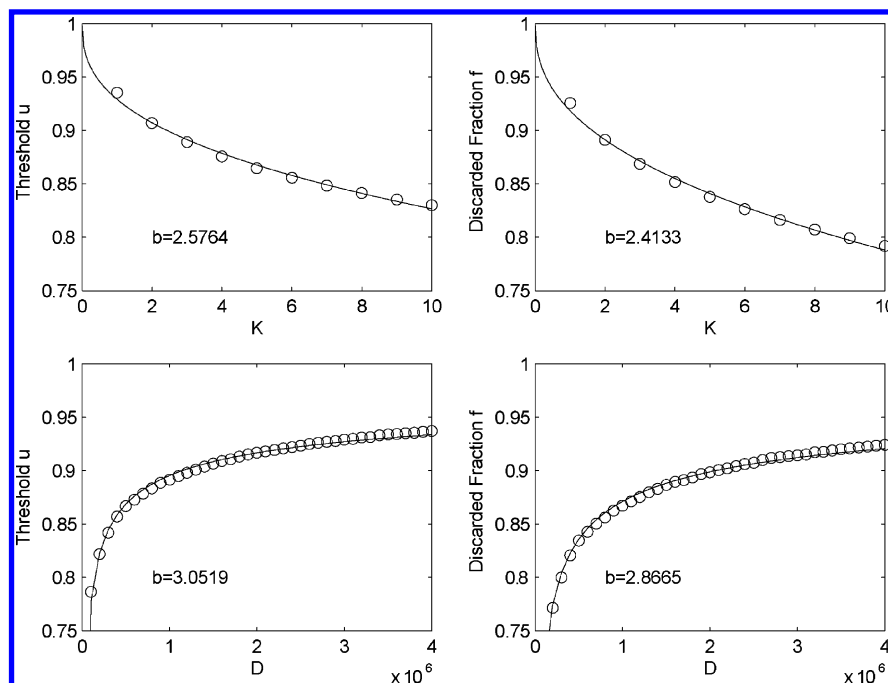


Figure 11. Plots (two upper) corresponding to an experiment where K is varied and $|D|$ is held constant at 4 099 792, where the results are averaged over 5000 separate queries randomly chosen from ChemDB, and plots (two lower) corresponding to an experiment where $|D|$ is varied and K is held constant at 1, where the data is averaged over 1000 separate queries randomly chosen from ChemDB. The lines are the best fit curves using the functional form given by $y = 1 - C(K/|D|)^{1/b}$, where b and C are the fit parameters and y corresponds either to u or to the fraction pruned from the database. This equation fits the data very closely with similar values for b . One can notice a small, but systematic, misfit between the empirical points and the theoretical curve $y = 1 - C(K/|D|)^{1/b}$. This can be entirely eliminated by introducing one additional offset parameter and fitting $y = 1 - [C_1 + C_2(K/|D|)^{1/b}]$ to the data.

from which we can derive

$$u \approx 1 - \left(\frac{bK}{\gamma|D|} \right)^{1/b} = 1 - \left[\frac{bK \Gamma(a) \Gamma(b)}{\Gamma(a+b) |D|} \right]^{1/b} \quad (41)$$

Thus, as a function of database size, the threshold u scales like $1 - C|D|^{-1/b}$. In turn, this threshold can be entered in the equations derived above, especially when the threshold approaches 1, to obtain the scaling of the speedup factor as a function of database size. For instance, by substituting this threshold in eq 31, we get

$$1 - f_D \approx \frac{\mu}{\sigma\sqrt{\pi}} \left[\frac{bK \Gamma(a) \Gamma(b)}{\Gamma(a+b) |D|} \right]^{1/b} \approx C|D|^{-1/b} \quad (42)$$

and similarly with eq 36, when the threshold approaches 1. The fraction of the database to be searched scales like $C|D|^{1-(1/b)}$. Thus, the fraction of the database to be searched becomes smaller and smaller as the size $|D|$ of the database increases and the speedup factor grows like $|D|^{1/b}$.

It is worth noting that more subtle analyses are possible by modeling f_D as a mixture of two (or more) β distributions where, for example, one component has a large mixing coefficient and corresponds to the overwhelming majority of molecules that do not match the query, while the other component has a small mixing coefficient and corresponds to the molecules that have a high similarity to the query. Since in the calculation above we are interested exclusively in the right tail of f_D near 1, the same analysis can be applied to these more complex models by focusing on the

corresponding β component and including the mixing coefficient in the proportionality constant.

Figure 11, drawn for fingerprints of length $N = 512$, shows the excellent agreement between the theoretical expression (eq 42) and the empirical values. With Tanimoto similarity, this yields values of b on the order of 2.5; thus, the fraction of the database to be searched scales like $O(|D|^{1-(1/b)}) = O(|D|^{0.6})$, yielding a sublinear search time.

While the simulation results show that these scaling formulas hold for current database sizes with millions of chemicals, it is clear that the scaling breaks down as $|D| \rightarrow \infty$ simply because of the discrete nature of the bins associated with the values of B . For extremely large values of $|D|$, the top K molecules \mathcal{B} that are most similar to a query molecule \mathcal{A} must be in the same bin as \mathcal{A} and must satisfy $B = A$. Therefore, once the database is large enough to reach this regime, by the current methods, one must search all the molecules contained in the same bin as \mathcal{A} and the speedup becomes linear again; that is, one needs to search $P_N(A)|D|$ records. It is possible to estimate the size $D^*(N)$ of the database at which this regime takes over. Clearly, D^* depends on the fingerprint length N and increases with N . To estimate D^* on the basis of an average case analysis, we notice that, when the linear regime is in effect, for a query \mathcal{A} , we need to search a fraction $P_N(A)$ of molecules; thus, on average, we need to search a fraction of the database equal to

$$\phi(N) = \sum_A P_N(A)^2 \quad (43)$$

and discard the fraction $1 - \phi$. We can then apply this value on the y axis of the curves in Figure 10 or 11 to find the




Figure 12. Plot of points, each representing a molecule in the random subset of 50 000 molecules from ChemDB, using binary fingerprints of length $N = 512$. A random jitter uniform over $[-0.5, 0.5]$ is injected in each coordinate to improve readability. The red curve corresponds to the predicted relationship between A and A^* using the asymptotic approximation to the binomial model (eq 78). The predicted values correspond closely to the expected empirical values.

Table 2. Actual Search-Time Benchmarks Obtained Searching the Entire ChemDB Database, with about 5M Compounds Using a 2.4 MHz AMD Opteron Processor with 2 GB of Memory^a

data set	size	time ($t = 0.9$)	$1 - f$	time ($K = 10$)	$1 - f$	time ($t = 0.9, K = 10$)	$1 - f$
Cox2	128	0.79	0.17	3.53	0.76	0.78	0.17
Estrogen	55	0.60	0.12	2.03	0.43	0.52	0.11
Gelatinase A	43	0.77	0.16	3.31	0.71	0.77	0.16
Neuraminidase	17	0.70	0.14	2.74	0.59	0.66	0.14
p38 MAP kinase	25	0.90	0.18	3.30	0.71	0.87	0.18
Thrombin	67	0.91	0.19	3.27	0.70	0.88	0.19
ChemDB Queries	1000	0.27	0.06	1.12	0.24	0.26	0.06
Random ChemDB	100	0.64	0.14	1.23	0.27	0.58	0.12

^a Searches are carried out using the Tanimoto similarity measure with threshold ($t = 0.9$) or top ten ($K = 10$) or both. Search times for the single-molecule query are expressed in seconds and are averaged over each data set. The data sets correspond to the six Stahl and Rarey¹⁴ data sets, a random set of 1000 queries extracted from the set of actual ChemDB queries, and a random set of 100 queries taken from ChemDB. The fraction of the database that needs to be searched is given by $1 - f$.

corresponding value of $|D|$. Using the curves and equations of Figure 11, this gives $\phi(N) = C[K/D^*(N)]^{1/b}$ or $\phi(N) = C_1 + C_2[K/D^*(N)]^{1/b}$ resulting in

$$D^*(N) \approx K \left[\frac{C}{\phi(N)} \right]^b \text{ or } D^*(N) \approx K \left[\frac{C_2}{\phi(N) - C_1} \right]^b \quad (44)$$

For $N = 512$ and the values in Figure 11, we get $\phi = 0.0032$ and $|D|/K = 6.7268 \times 10^7$, using the more accurate model with an offset. Thus, the equations predict that the speedup will transition from sublinear to linear when $D^*(N)$ is about 67 million compounds with $K = 1$ and 670 million compounds when $K = 10$. These estimates are sensitive to the values of C and b and therefore should only be considered as indicative. However, they clearly indicate that for the foreseeable future the speedup will remain sublinear, and

this is even more true with longer fingerprints, of length $N = 1024$ and beyond, which are widely used.

8. CONCLUSION

Current fingerprint search systems often require sequentially scanning the entire fingerprint database. While this search strategy works for small to moderately sized databases which can be stored in memory, it does not scale up well. In most practical applications, the user is not interested in computing all similarity values, but only the top hits, and perhaps estimating the histogram of the entire distribution of similarity scores. Here, we have shown how, using prestored information about the number of bits set to 1 in each binary fingerprint, or the total number of counts in a nonbinary fingerprint, we can greatly restrict the number of compounds for which a similarity score must be computed.

The method uses simple mathematical bounds on fingerprint similarity measures to search only a fraction of the database using the prestored information. With searches involving a fixed linear threshold, the method yields linear speedup where one needs to search only a fraction $(1 - f)$ of the database. We have shown how the factor $1 - f$ depends on threshold and fingerprint length and can lead to speedups of 1 order of magnitude or more. With searches aimed at retrieving the top K hits, the method yields even better sublinear speedups where, in a typical case, the fraction of the database to be searched scales like $|D|^{0.6}$. The method requires no tuning and remains exact in the sense that it returns the same results as if the entire database was searched. As shown in the benchmark experiments reported in Table 2, when applied to the 5M compounds in ChemDB, the method delivers subsecond search times on a single desktop machine.

ACKNOWLEDGMENT

Work supported by an NIH Biomedical Informatics Training Grant (LM-07443-01) and an NSF MRI Grant (EIA-0321390) to P.B., by the UCI Medical Scientist Training Program, and by a Harvey Fellowship to S.J.S. We would like also to acknowledge the OpenBabel project and OpenEye Scientific Software for their free software academic license and Drs. Chamberlin, Nowick, and Weiss for their useful feedback.

APPENDIX A

Extensions to Other Similarity Measures. Fingerprint similarity can be calculated using many different measures. Holliday et al.¹⁸ compare a comprehensive list of fingerprint similarities and distances. Using their nomenclature, we show here how to calculate bounds on these measures, essentially by expressing each measure in terms of unions and intersections, studying the corresponding derivatives, and applying corresponding bounds.

As previously derived, the *Jaccard/Tanimoto* measure is defined as, and bounded by,

$$S(\vec{A}, \vec{B}) = \frac{A \cap B}{A \cup B} \leq T(A, B) = \frac{\min(A, B)}{\max(A, B)} \quad (45)$$

The *Baroni-Urbani/Buser* measure is similar to the Jaccard/Tanimoto measure and is defined as

$$S(\vec{A}, \vec{B}) = \frac{\sqrt{(A \cap B)(N - A \cup B)} + A \cap B}{\sqrt{(A \cap B)(N - A \cup B)} + A \cup B} \quad (46)$$

and using the identity $\min(A, B) \max(A, B) = AB$, we can derive the bound

$$S(\vec{A}, \vec{B}) \leq T(A, B) = \frac{\sqrt{N \min(A, B) - AB} + \min(A, B)}{\sqrt{N \min(A, B) - AB} + \max(A, B)} \quad (47)$$

Closely related to the Jaccard/Tanimoto measure, the *Ochiai/Cosine* measure is defined as, and bounded by,

$$S(\vec{A}, \vec{B}) = \frac{A \cap B}{\sqrt{AB}} \leq T(A, B) = \frac{\min(A, B)}{\sqrt{AB}} \quad (48)$$

Also closely related to the Jaccard/Tanimoto measure, the *Pearson* measure is defined as

$$S(\vec{A}, \vec{B}) = \frac{N(A \cap B) - AB}{\sqrt{AB(N - A)(N - B)}} \quad (49)$$

and bounded by

$$S(\vec{A}, \vec{B}) \leq T(A, B) = \frac{N \min(A, B) - AB}{\sqrt{AB(N - A)(N - B)}} \quad (50)$$

The rather complex *Stiles* measure produces nearly identical rankings as those of the Pearson measure. It is defined as

$$S(\vec{A}, \vec{B}) = \log_{10} \frac{N \left(\left| N(A \cap B) - AB \right| - \frac{N}{2} \right)^2}{AB(N - A)(N - B)} \quad (51)$$

and bounded by

$$S(\vec{A}, \vec{B}) \leq T(A, B) = \log_{10} \frac{N \left[N \min(A, B) - AB - \frac{N}{2} \right]^2}{AB(N - A)(N - B)} \quad (52)$$

The *McConnaughey* and *Dennis* measures also are both closely correlated with the Jaccard/Tanimoto measure. The *McConnaughey* measure is defined as

$$S(\vec{A}, \vec{B}) = \frac{(A \cap B)(2 - A - B) + AB}{AB} \quad (53)$$

and bounded by

$$S(\vec{A}, \vec{B}) \leq T(A, B) = \frac{\min(A, B)(2 - A - B) + AB}{AB} \quad (54)$$

While the *Dennis* measure is defined as, and bounded by,

$$S(\vec{A}, \vec{B}) = \frac{N(A \cap B) - AB}{\sqrt{NAB}} \leq T(A, B) = \frac{N \min(A, B) - AB}{\sqrt{NAB}} \quad (55)$$

Russel/Rao, perhaps the simplest measure, performs surprisingly well: in some tests better than the more commonly used Jaccard/Tanimoto measure. It is defined as, and bounded by,

$$S(\vec{A}, \vec{B}) = \frac{A \cap B}{N} \leq T(A, B) = \frac{\min(A, B)}{N} \quad (56)$$

The *Dice* and *Sokal/Sneath(1)* measures are equivalent to the Tversky similarity measure with $\alpha = \beta = 0.5$ and $\alpha = \beta = 2$, respectively. Noting this relationship, we can use the

bounds on the Tversky similarity to derive bounds on these measures. For example, for the Dice measure we have

$$S(\vec{A}, \vec{B}) = \frac{2(A \cap B)}{A + B} \leq T(A, B) = \frac{2 \min(A, B)}{A + B} \quad (57)$$

The *Sokal/Sneath(2)* measure is defined as

$$S(\vec{A}, \vec{B}) = \frac{2[N - A - B + 2(A \cap B)]}{2N - A - B + 2(A \cap B)} \quad (58)$$

Using the identity $A + B - 2 \min(A, B) = |A - B|$, we can derive its bound

$$S(\vec{A}, \vec{B}) \leq T(A, B) = \frac{2N - 2|A - B|}{2N - |A - B|} \quad (59)$$

We can use the same identity to define and bound the *Sokal/Sneath(3)* measure

$$S(\vec{A}, \vec{B}) = \frac{N - A - B + 2(A \cap B)}{A + B - 2(A \cap B)} \leq T(A, B) = \frac{N - |A - B|}{|A - B|} \quad (60)$$

Similarly, the *Simple Matching* measure, which can also be viewed as the complement of the Mean Manhattan Distance, is defined as, and bounded by,

$$S(\vec{A}, \vec{B}) = \frac{N - A - B + 2(A \cap B)}{N} \leq T(A, B) = \frac{N - |A - B|}{N} \quad (61)$$

Similarly, the *Kulczynski(1)* measure is defined as, and bounded by,

$$S(\vec{A}, \vec{B}) = \frac{A \cap B}{A + B - 2(A \cap B)} \leq T(A, B) = \frac{\min(A, B)}{|A - B|} \quad (62)$$

When the identity $\min(A, B)/(AB) = 1/\max(A, B)$ is again used, the *Kulczynski(2)* measure is defined as, and bounded by,

$$S(\vec{A}, \vec{B}) = \frac{(A \cap B)(A + B)}{2AB} \leq T(A, B) = \frac{A + B}{2 \max(A, B)} \quad (63)$$

Similarly, the *Hamann* measure is defined as, and bounded by,

$$S(\vec{A}, \vec{B}) = \frac{N + 4(A \cap B) - 2A - 2B}{N} \leq T(A, B) = \frac{N - 2|A - B|}{N} \quad (64)$$

The *Rogers/Tanimoto* measure is similar to the *Sokal/Sneath(3)* measure. It is defined as, and bounded by,

$$S(\vec{A}, \vec{B}) = \frac{N - A - B + 2(A \cap B)}{N + A + B - 2(A \cap B)} \leq T(A, B) = \frac{N - |A - B|}{N + |A - B|} \quad (65)$$

The *Forbes* measure is defined as, and bounded by,

$$S(\vec{A}, \vec{B}) = \frac{N(A \cap B)}{AB} \leq T(A, B) = \frac{N}{\max(A, B)} \quad (66)$$

Closely correlated with the *Forbes* measure, the *Yule* measure is defined as

$$S(\vec{A}, \vec{B}) = \frac{N(A \cap B) - AB}{(A \cap B)(N - 2A - 2B + 2) + AB} \quad (67)$$

and bounded by

$$S(\vec{A}, \vec{B}) \leq T(A, B) = \frac{N \min(A, B) - AB}{\min(A, B)(N - 2A - 2B + 2) + AB} = \frac{N - \max(A, B)}{N - 2A - 2B + 2 + \max(A, B)} \quad (68)$$

The *Fossum* measure is defined and bounded by

$$S(\vec{A}, \vec{B}) = \frac{N[(A \cap B) - 0.5]^2}{AB} \leq T(A, B) = \frac{N[\min(A, B) - 0.5]^2}{AB} \quad (69)$$

The *Simpson* measure is defined as

$$S(\vec{A}, \vec{B}) = \frac{A \cap B}{\min(A, B)} \quad (70)$$

In this case, the obvious upper bound is 1, and it is achieved when $A \neq 0$ and $B \neq 0$.

$$S(\vec{A}, \vec{B}) \leq T(A, B) = \frac{\min(A, B)}{\min(A, B)} = 1 \quad (71)$$

This upper bound does not lead to any search space reduction.

In some applications, such as selecting a diverse data set, one is interested in minimizing similarity (or maximizing distance) instead. This approach is meaningful only with certain metrics. To maximize the distance between fingerprints, we must minimize the overlap between \vec{A} and \vec{B} , and so we can lower bound $(A \cap B)$ with $\max(A + B - N, 0)$, the minimum possible overlap. For example, this can be applied to the *Mean Manhattan Distance*, which is also the L1 norm or the city-block distance. It is the complement of the *Simple Matching* measure, and in the binary case, it is rank-equivalent to the Euclidian distance. It is defined as, and bounded by,

$$S(\vec{A}, \vec{B}) = \frac{A + B - 2(A \cap B)}{N} \leq T(A, B) = \frac{A + B - 2 \max(A + B - N, 0)}{N} \quad (72)$$

Similarly, the *Normalized Euclidian* distance is the familiar L2 norm, defined as

$$S(\vec{A}, \vec{B}) = \sqrt{\frac{A + B - 2(A \cap B)}{N}} \quad (73)$$

It can be bounded in the same way by

$$S(\vec{A}, \vec{B}) \leq T(A, B) = \sqrt{\frac{A + B - 2 \max(A + B - N, 0)}{N}} \quad (74)$$

In short, we have shown how the same principles can be applied essentially to all the other fingerprint similarity/distance measures that are found in the literature.

APPENDIX B

Estimation of A Given A_* . To estimate A from A_* , we assume that the order of the bits in the long uncompressed fingerprints has been randomized, so that we can use a simple binomial approximation where we assume that the bits in the uncompressed fingerprints are set to 1 according to a binomial coin-flip process $\mathcal{B}(N_*, \alpha)$ with probability $\alpha = A_*/N_*$. Then, with modulo compression, the probability of setting a given bit to 0 in the compressed A is $(1 - \alpha)^k$, where $k = N_*/N$. Therefore, the corresponding distribution for A is binomial $\mathcal{B}(N, p)$ with

$$P(A|\alpha) = \mathcal{B}(n, p) \text{ with } p = 1 - (1 - \alpha)^k \quad (75)$$

Therefore, given α (or A_*)

$$E(A|\alpha) = N[1 - (1 - \alpha)^k] \quad (76)$$

Thus, given A_* , we can estimate A by

$$E(A|A_*) \approx N \left[1 - \left(1 - \frac{A_*}{N_*} \right)^{N_*/N} \right] \quad (77)$$

Asymptotically, for a large N_* , we can use

$$E(A|A_*) \approx N(1 - e^{-A_*/N}) \quad (78)$$

Equations 77 and 78 provide very good approximations to the true value, as shown in Figure 12.

REFERENCES AND NOTES

- (1) Irwin, J. J.; Shoichet, B. K. ZINC—A Free Database of Commercially Available Compounds for Virtual Screening. *J. Chem. Inf. Model.* **2005**, *45*, 177–182.
- (2) Chen, J.; Swamidass, S. J.; Dou, Y.; Baldi, P. ChemDB: A Public Database of Small Molecules and Related Chemoinformatics Resources. *Bioinformatics* **2005**, *21*, 4133–4139.
- (3) Fligner, M. A.; Verducci, J. S.; Blower, P. E. A Modification of the Jaccard/Tanimoto Similarity Index for Diverse Selection of Chemical Compounds Using Binary Strings. *Technometrics* **2002**, *44*, 110–119.
- (4) Flower, D. R. On the Properties of Bit String-Based Measures of Chemical Similarity. *J. Chem. Inf. Comput. Sci.* **1998**, *38*, 378–386.
- (5) James, C. A.; Weininger, D.; Delany, J. *Daylight Theory Manual*; Daylight Chemical Information Systems, Inc.: Aliso Viejo, CA, 2004. Available at <http://www.daylight.com/dayhtml/doc/theory/theory-toc.html> (accessed Jan 2007).
- (6) Xue, L.; Godden, J. F.; Stahura, F. L.; Bajorath, J. Profile Scaling Increases the Similarity Search Performance of Molecular Fingerprints Containing Numerical Descriptors and Structural Keys. *J. Chem. Inf. Comput. Sci.* **2003**, *43*, 1218–1225.
- (7) Xue, L.; Stahura, F. L.; Bajorath, J. Similarity Search Profiling Reveals Effects of Fingerprint Scaling in Virtual Screening. *J. Chem. Inf. Comput. Sci.* **2004**, *44*, 2032–2039.
- (8) Bohacek, R. S.; McMartin, C.; Guida, W. C. The Art and Practice of Structure-Based Drug Design: A Molecular Modelling Perspective. *Med. Res. Rev.* **1996**, *16*, 3–50.
- (9) Tversky, A. Features of Similarity. *Psychol. Rev.* **1977**, *84*, 327–352.
- (10) Rouvray, D. Definition and Role of Similarity Concepts in the Chemical and Physical Sciences. *J. Chem. Inf. Comput. Sci.* **1992**, *32*, 580–586.
- (11) Swamidass, S. J.; Chen, J.; Bruand, J.; Phung, P.; Ralaivola, L.; Baldi, P. Kernels for Small Molecules and the Prediction of Mutagenicity, Toxicity, and Anti-Cancer Activity. *Bioinformatics* **2005**, *21*, i359–368.
- (12) van Rijsbergen, C. J. *Information Retrieval*; Butterworths: London, U. K., 1978.
- (13) Ralaivola, L.; Swamidass, S. J.; Saigo, H.; Baldi, P. Graph Kernels for Chemical Informatics. *Neural Networks* **2005**, *18*, 1093–1110.
- (14) Stahl, M.; Rarey, M. Detailed Analysis of Scoring Functions for Virtual Screening. *J. Med. Chem.* **2001**, *44*, 1035–1042.
- (15) Downs, G. M.; Gillet, V. J.; Holliday, J. D.; Lynch, M. F. Computer Storage and Retrieval of Generic Chemical Structures in Patents. 10. The Generation and Logical Bubble-Up of Ring Screens for Structurally-Explicit Generics. *J. Chem. Inf. Comput. Sci.* **1989**, *29*, 215–224.
- (16) Shemetulskis, N. E.; Weininger, D.; Blankley, C. J.; Yang, J. J.; Humblet, C. Stigmata: An Algorithm to Determine Structural Commonalities in Diverse Datasets. *J. Chem. Inf. Comput. Sci.* **1996**, *36*, 862–871.
- (17) Wild, D. J.; Blankley, C. J. VisualSAR: A Web-Based Application for Clustering, Structure Browsing, and Structure–Activity Relationship Study. *J. Mol. Graphics Modell.* **1999**, *17*, 85–89; 120–125.
- (18) Holliday, J. D.; Hu, C. Y.; Willett, P. Grouping of Coefficients for the Calculation of Inter-Molecular Similarity and Dissimilarity Using 2d Fragment Bit-Strings. *Comb. Chem. High Throughput Screen.* **2002**, *5*, 155–166.

CI600358F